

logitSVD

Combines the advantages of classic logit models with the SVD approach for building product recommendation systems




Predicts the probability to buy a product or to rate the product e.g. with a certain number of stars



Strong predictive power with a minimal number of parameter to avoid overfitting



Avoids cold start problem through the use of user features



Selection of features is supported by a feature importance measure (Wald test)



Highly transparent model with interpretable parameters



Building separate models for each product is a special case of logitSVD

A feature based SVD for binary user-item-matrices

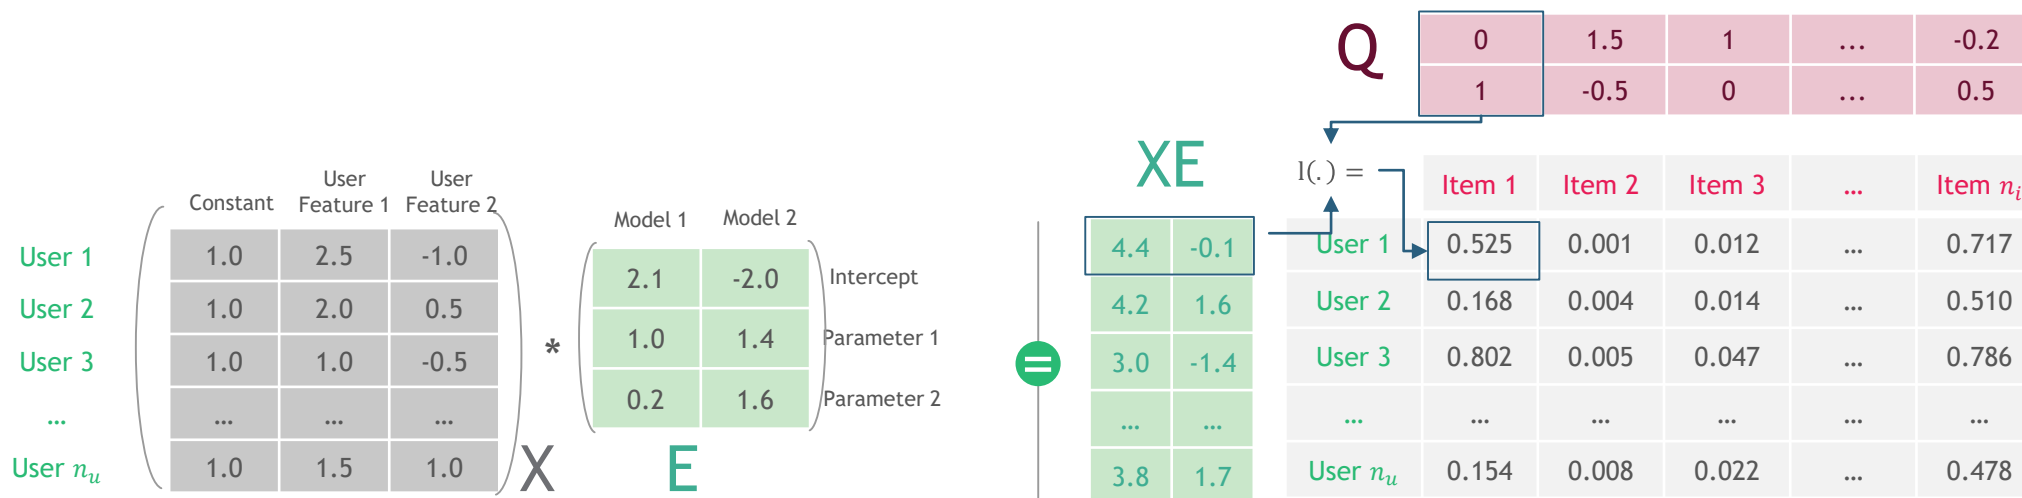
Definition¹

Let $R \in \{0,1,NaN\}^{n_u \times n_i}$ be a binary matrix with entries 0, 1 or missing value for n_u users and n_i items and let n_f features $X \in \mathbb{R}^{n_u \times n_f}$ for each user u be given. The probability $p_{u,i} = P(r_{u,i} = 1)$ that $r_{u,i} = 1$ can then be approximated by the binary logit SVD:

$$S = X E Q, \quad p_{u,i} = l(s_{u,i}) = \frac{1}{1 + e^{-s_{u,i}}}$$

with parameter matrices $E \in \mathbb{R}^{n_f \times n_d}$ and $Q \in \mathbb{R}^{n_d \times n_i}$. n_d is a design parameter, which can be interpreted as the number of estimated models.

Illustration



1. We are using the matrix notation $A = (a_{i,j})_{i \leq n, j \leq m}$

Newton's method boosted maximum likelihood parameter estimation for binary user-item-matrices

Maximum likelihood

The parameter matrixes E and Q are estimated such the log-likelihood to observe $R \in \{0,1,NaN\}^{n_u \times n_i}$ is maximized

$$\log L(E, Q) = \sum_{u \leq n_u, i \leq n_i, r_{u,i} \neq NaN} \left[r_{u,i} \log(l(s_{u,i})) + (1 - r_{u,i}) \log(1 - l(s_{u,i})) \right]$$

with $S = X E Q \in \mathbb{R}^{n_u \times n_i}$ and $l(x) = \frac{1}{1+e^{-x}}$.

Note that only known product usages, i.e. $r_{u,i} \neq NaN$, are considered in the likelihood.

Full Newton's and alternating Newton's method

The logitVD package offers two options for solving the non-linear likelihood optimization problem:

- straight Newton's method with line search to simultaneously calculate E and Q and
- alternating Newton's methods with line search to iteratively optimize E with fixed Q and subsequently optimize Q with fixed E

For most problems the alternating approach is quicker and more stable. It can be shown that each of the alternating optimization problems (i.e. for E with fixed Q and Q with fixed E) is convex with s.p.d. Hessian.

A feature based SVD for user-item-matrices with several ordered classes

Definition¹



Let $R \in \{0, 1, \dots, n_k, NaN\}^{n_u \times n_i}$ be the user-item-matrix describing the rating given by user u to item i . The ratings $0 \leq r_{u,i} \leq n_k$ are assumed to be ordered, i.e. all ratings $r_{u,i}$ with $r_{u,i} < r_{v,j}$ are better/worse than the ratings $r_{v,j}$. The probability $p_{u,i}^{(<k)} = P(r_{u,i} < k)$ that the rating $r_{u,i}$ is smaller than k can then be approximated by the feature based SVD:

$$S = X E Q, \quad p_{u,i}^{(<k)} = l(s_{u,i} - t_k) = \frac{1}{1 + e^{-s_{u,i} - t_k}}$$

with parameter matrices $E \in \mathbb{R}^{n_f \times n_d}$ and $Q \in \mathbb{R}^{n_d \times n_i}$ and intercepts $t_k \in \mathbb{R}$ for $1 \leq k \leq n_k$ and $t_0 = -\infty$, $t_{n_k+1} = \infty$.

The probability $p_{u,i}^{(k)}$ that item i is rated k by user u rating $r_{u,i}$ is then $p_{u,i}^{(k)} = P(r_{u,i} = k) = p_{u,i}^{(<k+1)} - p_{u,i}^{(<k)}$

Illustration



1. We are using the matrix notation $A = (a_{i,j})_{i \leq n, j \leq m}$

Newton's method boosted maximum likelihood parameter estimation for ordered target classes

Maximum likelihood

The parameter matrixes E , Q and the intercepts t are estimated such the log-likelihood to observe $R \in \{0, 1, \dots, n_k, NaN\}^{n_u \times n_i}$ is maximized

$$\log L(E, Q, t) = \sum_{k \leq n_k} \left\{ \sum_{u \leq n_u, i \leq n_i, r_{u,i}=k} \left[\log(l(s_{u,i} + t_{k+1})) - \log(l(s_{u,i} + t_k)) \right] \right\}$$

with $S = X E Q \in \mathbb{R}^{n_u \times n_i}$ and $l(x) = \frac{1}{1+e^{-x}}$.

Note that only known product usages, i.e. $r_{u,i} \neq NaN$, are considered in the likelihood.

Two different alternating Newton's method

The logitSVD package offers two options for solving the non-linear likelihood optimization problem:

- an alternating Newton's method with line search to iteratively optimize simultaneously t and E for fixed Q and then t and Q for fixed E
- An alternating Newton's methods with line search to iteratively optimize all 3 elements separately, i.e. first E with fixed t and Q , then t for fixed E and Q , and subsequently Q with fixed t and E .

For most problems, both approaches are successful. Typically, the first approach is quicker but less stable, i.e. one of the Hessians might not be positive definite.

logitSVD is implemented in a Python package (1/2)

Function call

```
P, C, Z, E, Q, t, z_score, p_value = logitSVD(X, R, depth, la, E = None, Q = None, t=None, method="alternating", tol = 1e-4, maxit = 20, tolNewton = None, maxitNewton = 100, verbose = "warn")
```

Parameter

```
X      : ndarray[nuser,nfeature], user feature vectors
R      : ndarray[nuser,nitem], user-item-matrix (target)
depth  : int, model parameter, depth of the embeddings = number of different models
la     : float, regularization paramter
E      : ndarray[nfeature,depth], initial solution for the feature weights (embeddings)
Q      : ndarray[depth,nitem], initial solution for the item embeddings (model combination parameter)
t      : ndarray[max(R)], initial solution for intercepts (only for multinomial case)
method : string, binary: alternating [alter], fullNewton [full], alter_full, i.e. first alter, then
        fullNewton
        multinomial: alternating2 [alter2], 2 alternating steps 1. Q,t and 2. E,t
        alternating3 [alter3], 3 alternating steps 1. Q, 2. t, 3. E
tol     : float, alternating methods stop if the reduction of the log-likelihood is smaller than tol
maxit   : int, maximum number of iterations of the alternating methods
tolNewton: float, Newton's method stops if the 2-norm of the gradient becomes smaller than tolNewton
maxitNewton: int, maximum number of iterations of Newton's method
verbose: string, ("none" | "warn" | "all"), print warnings and convergence progress. Default is "warn"
```

logitSVD is implemented in a Python package (2/2)

Function call

```
P, C, Z, E, Q, t, z_score, p_value = logitSVD(X, R, depth, la, E = None, Q = None, t=None, method="alternating", tol = 1e-4, maxit = 20, tolNewton = None, maxitNewton = 100, verbose = "warn")
```

Output

```
P      : binary: ndarray[nuser,nitem], P[u,i] is the probability that R[u,i] = 1  
        multinomial: ndarray[max(R),nuser,nitem], P[k,u,i] is the probability that R[u,i] = k  
C      : ndarray[nuser,nitem], most likely class - None for binary  
Z      : ndarray[nuser,nitem], Z = X E Q  
E      : ndarray[nfeature,depth], solution for the parameter E  
Q      : ndarray[depth,nitem], solution for the parameter Q  
t      : ndarray[max(R)], solution for the intercepts t - None for binary  
z-score: ndarray[nfeature,depth], z-score from Wald test for the parameter E  
p-value: ndarray[nfeature,depth], p-values from Wald test for the parameter E
```